

Chapter 2: Bits and Bytes

Topics

- Why bits?
- Binary Logic
- Representing information as bits
 - Binary/Hexadecimal
 - Byte representations
 - » ints, floats, double precision

Announcements

Moodle is having difficulty emailing announcements

Data Lab is due Friday Feb 3 by 11:55 pm

- Bit manipulation operations

C assessment quiz due Friday Jan 20 by noon

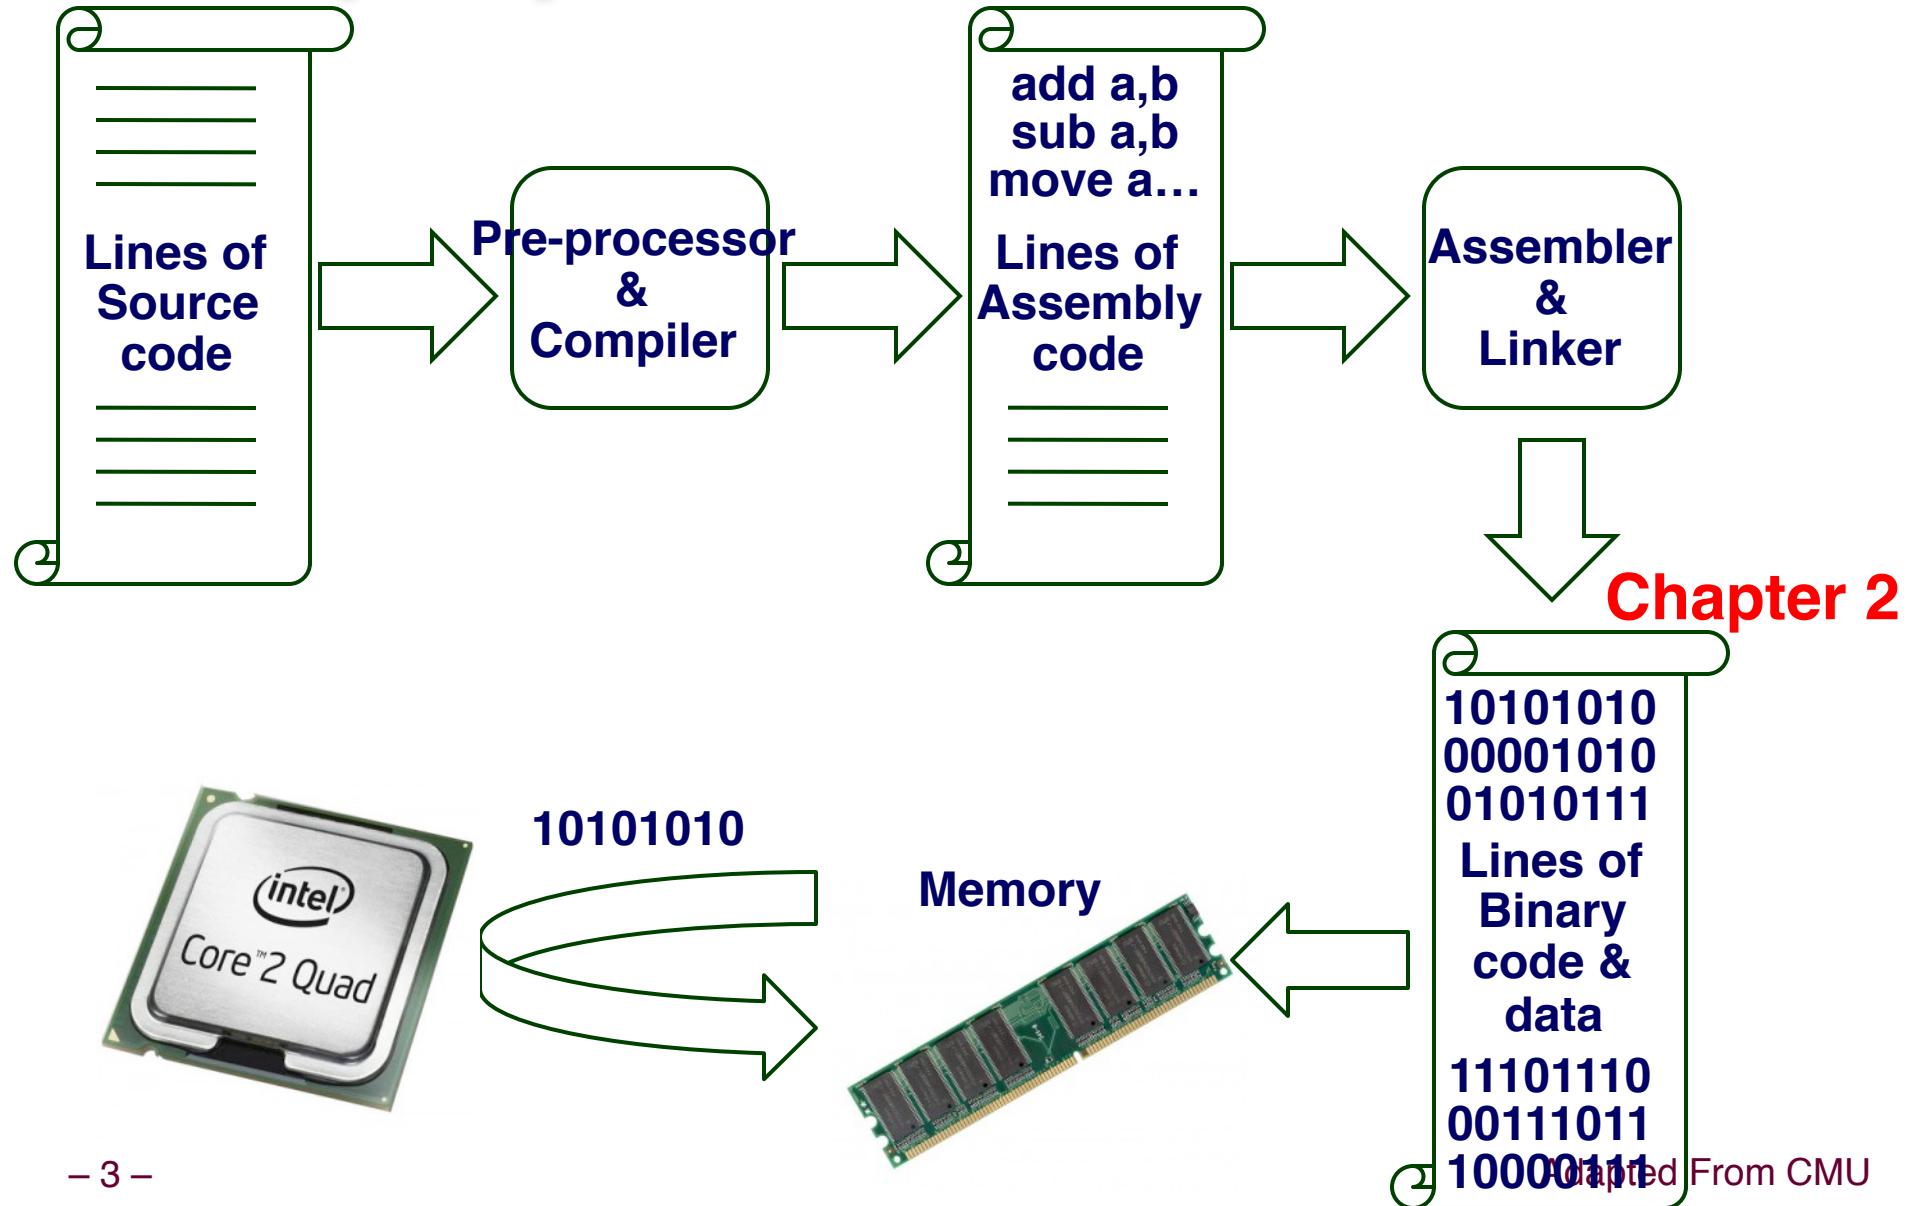
- Unlimited # of attempts
- If you score $\leq 70\%$, then attend C tutorial Friday evening (TA Yogesh will announce)

Chap 2 Data Quiz is due Monday Jan 30 by noon

- Addition, subtraction, signed, overflow

Read Chapter 2.1-2.3 and do practice problems

Recap: Systems In a Nutshell



Binary Representation – Why Bits?

$$\begin{aligned}0101_2 &= 0*2^3 + 1*2^2 + 0*2^1 + 1*2^0 \\&= 0 + 4 + 0 + 1 = 5_{10}\end{aligned}$$

Each binary digit is called a ‘bit’

Base 2 or Binary is easier to represent via electronic circuits than Base 10

- Digital transistor circuits can easily represent a ‘0’ and a ‘1’
 - A ‘1’ = +5 volts, a ‘0’ = 0 volts, so only 2 voltage levels
 - Don’t have to implement 10 different voltage levels as in base 10
- Easy to implement Base 2 arithmetic
 - much of Base 10 arithmetic using Base 2 arithmetic (integers are exact, but floats are approximations)
- Easy to implement digital logic (AND, OR, etc.) with binary

Common Digital Logic Operations

And “&”

- $A \& B = 1$ only when both $A=1$ and $B=1$

		B
&		0 1
A	0	0 0
	1	0 1

Or “|”

- $A | B = 1$ when either $A=1$ or $B=1$

		0 1
		0 1
0	0	0 1
	1	1 1

Not “~”

- $\sim A = 1$ only when $A=0$

		~
		1
0	0	1
	1	0

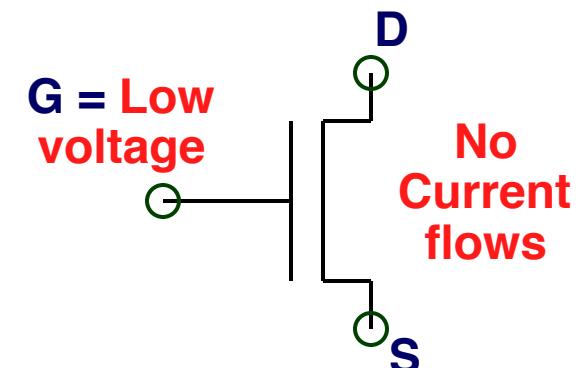
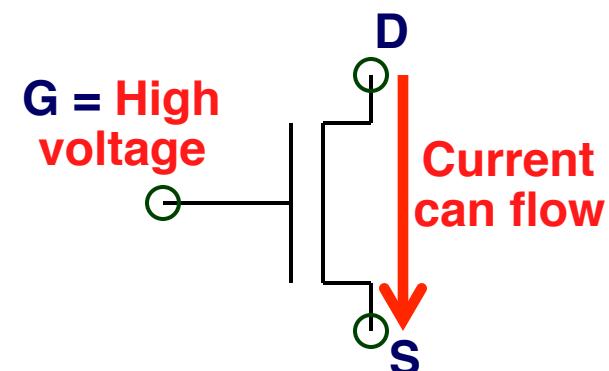
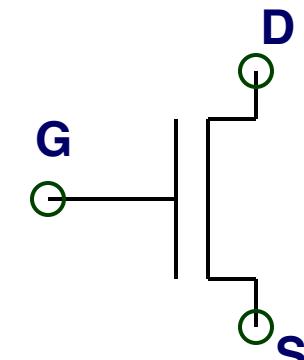
Exclusive-Or (Xor) “^”

- $A ^ B = 1$ when either $A=1$ or $B=1$, but not both

		0 1
^		0 1
0	0	0 1
	1	1 0

Implementing Digital Logic

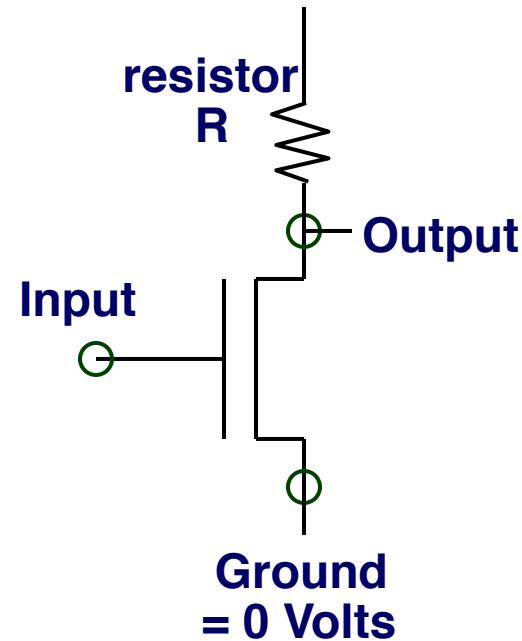
- **Silicon transistor (NMOS)**
 - Three terminals. Input is the Gate G
 - Acts like a simple ON/OFF switch
 - Current flows from high voltage to low voltage
 - When a high voltage is applied to the input Gate G, then this connects the two terminals D and S, i.e. closes the switch, and current flows between them
 - When a low voltage is applied to Gate, then the two terminals are disconnected, i.e. opens the switch, and no current flows



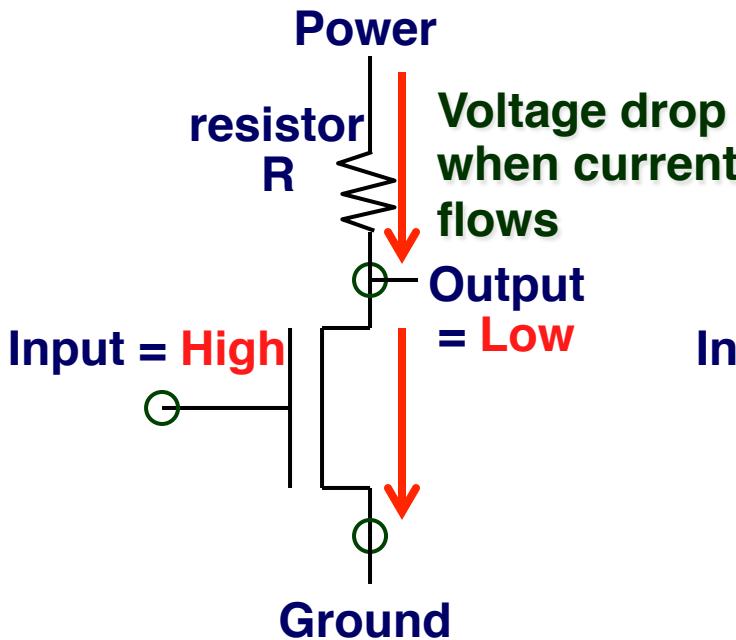
Implementing an Inverter or Not “~” Operation

An Inverter circuit

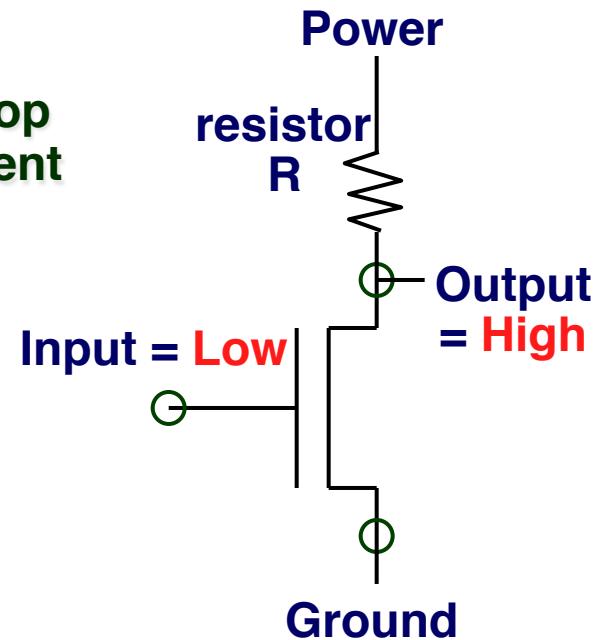
Power = +3 or +5 Volts



High input connects Output to Ground



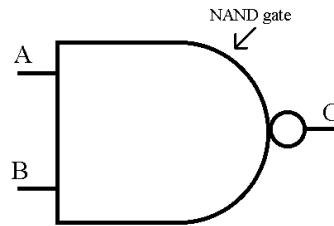
Low input means no current flows, so output stay high



- Thus we have an inverter or Not “~” operator:
 - Input ‘1’ (High voltage) => output ‘0’ (Low voltage)
 - Input ‘0’ (Low voltage) => output ‘1’ (High voltage)

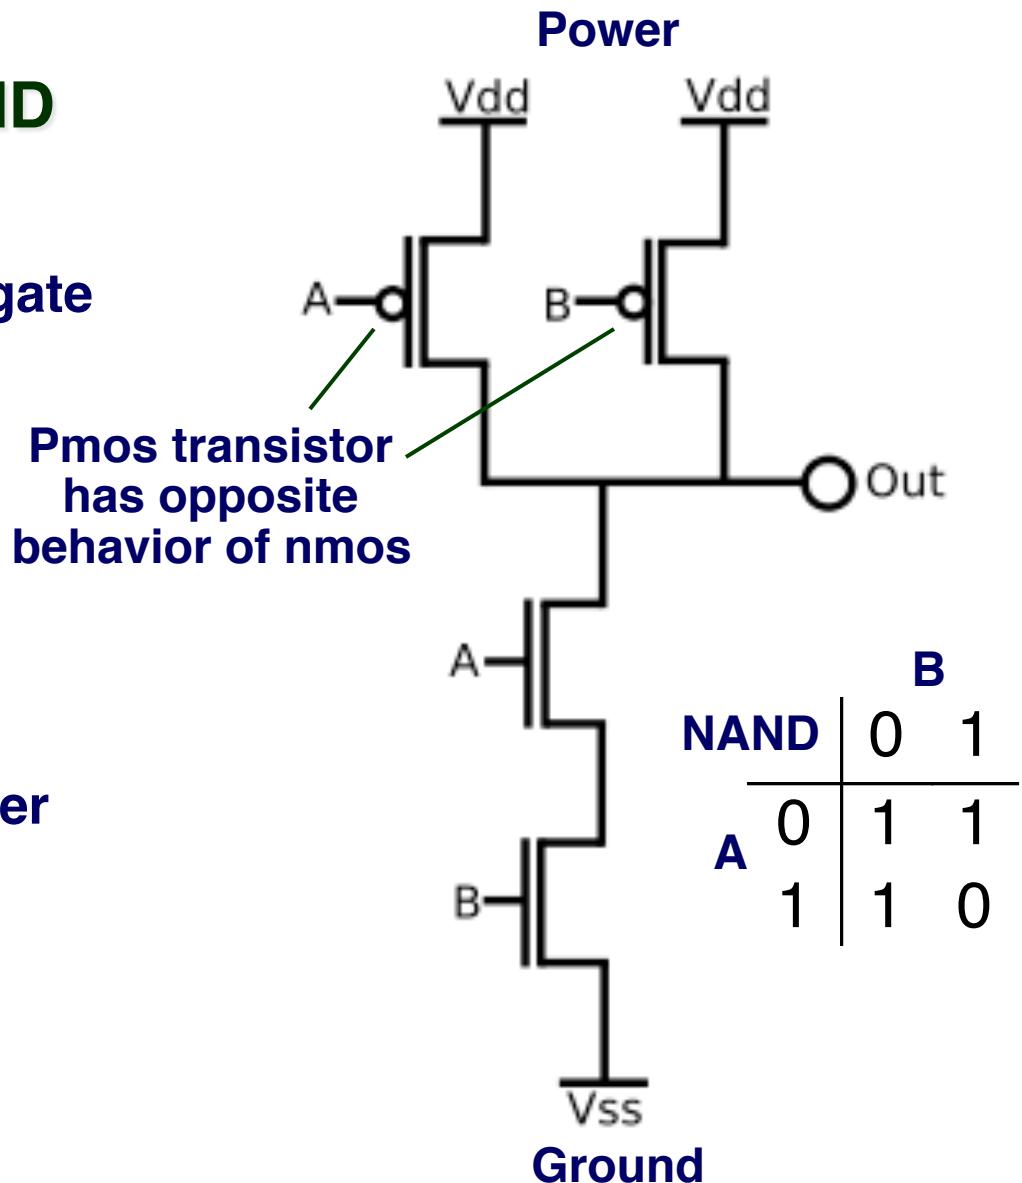
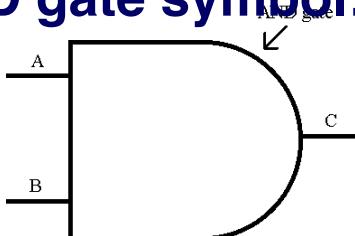
Implementing a NAND operation

- **NAND = Negative AND**
 - Output = $\sim(A \text{ AND } B)$
 - Symbol for a NAND gate is:

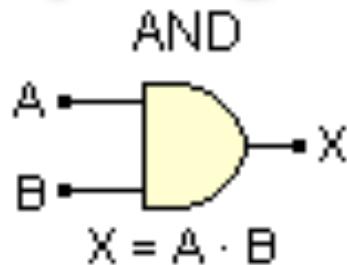


How do I make an AND gate?

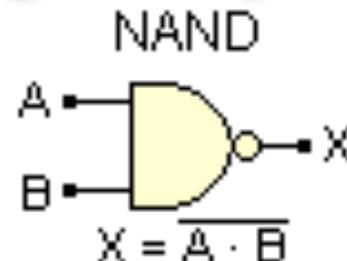
- Concatenate an inverter after a NAND gate!
- AND gate symbol:



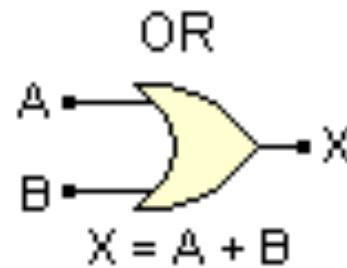
Binary Digital Logic Symbols



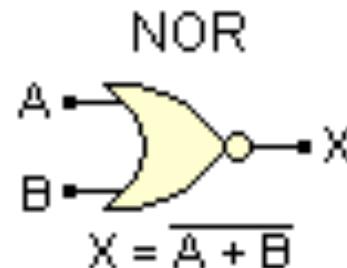
A	B	X
0	0	0
0	1	0
1	0	0
1	1	1



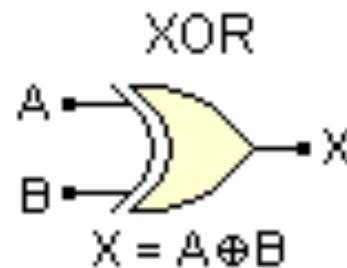
A	B	X
0	0	1
0	1	1
1	0	1
1	1	0



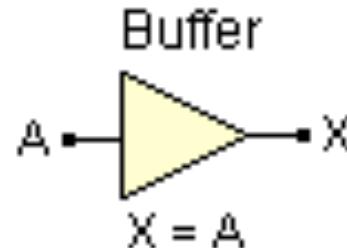
A	B	X
0	0	0
0	1	1
1	0	1
1	1	1



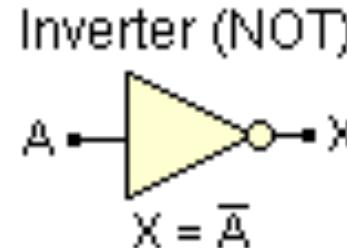
A	B	X
0	0	1
0	1	0
1	0	0
1	1	0



A	B	X
0	0	0
0	1	1
1	0	1
1	1	0



A	X
0	0
1	1



A	X
0	1
1	0

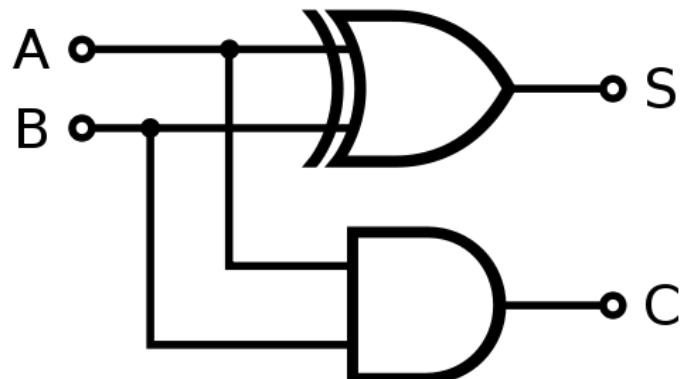
Building a Binary Adder

One-bit Half Adder circuit

- Sum bit $S = A \text{ XOR } B$
- Carry bit $C = A \text{ AND } B$
- One-bit Full Adder circuit extends Half Adder to also account for a bit carried in, not just the bit carried out
- Example: an 8-bit Full Adder can be built by concatenating 8 one-bit Full Adders

Bit A	Bit B	Sum S	Carry C
0	0	0	0
0	1	1	0
1	0	1	0
1	1	0	1

One-bit Half Adder Circuit



Binary Representation (2)

Base 2 vs. Decimal conversion example:

89 base 10 = 89_{10} = what in base 2 or binary?

 Find the largest power of 2 that is \leq # and then subtract.
Keep iterating until # = 0.

$$\begin{aligned} 89_{10} &= 64 + 16 + 8 + 1 \\ &= 2^6 + 2^4 + 2^3 + 2^0 \\ &= 0*2^7 + 1*2^6 + 0*2^5 + 1*2^4 + 1*2^3 + 0*2^2 + 0*2^1 + 1*2^0 \\ &= 0\ 1\ 0\ 1\ 1\ 0\ 0\ 1_2 \end{aligned}$$

Hexadecimal Representation

**Hexadecimal (base 16) is more convenient
(it's power of 2) and compact notation
for larger binary #'s**

- Consider a 32-bit number:

$10101110001100101000101000011001_2$

This is hard to read. So form 4-bit groups.

= $1010\ 1110\ 0011\ 0010\ 1000\ 1010\ 0001\ 1001_2$

We then use hexadecimal notation:

Use characters '0' to '9' and 'A' to 'F'

= A E 3 2 8 A 1 9₁₆

= 0xAE328A19

“0x” signifies hexadecimal or just hex.

Note lower case 0xae328a19 is also valid.

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

How We Use These Representations

Since these are just numbers, we can use them to describe anything with numerical value

We'll use them to describe the numerical *values* of data stored at various memory locations

- Different data values, like real and integer values, positive and negative

We'll also use them to describe the *addresses* of memory locations

Later, our code instructions will be encoded in binary as well

Byte-Addressed Memory

Byte = 8 bits

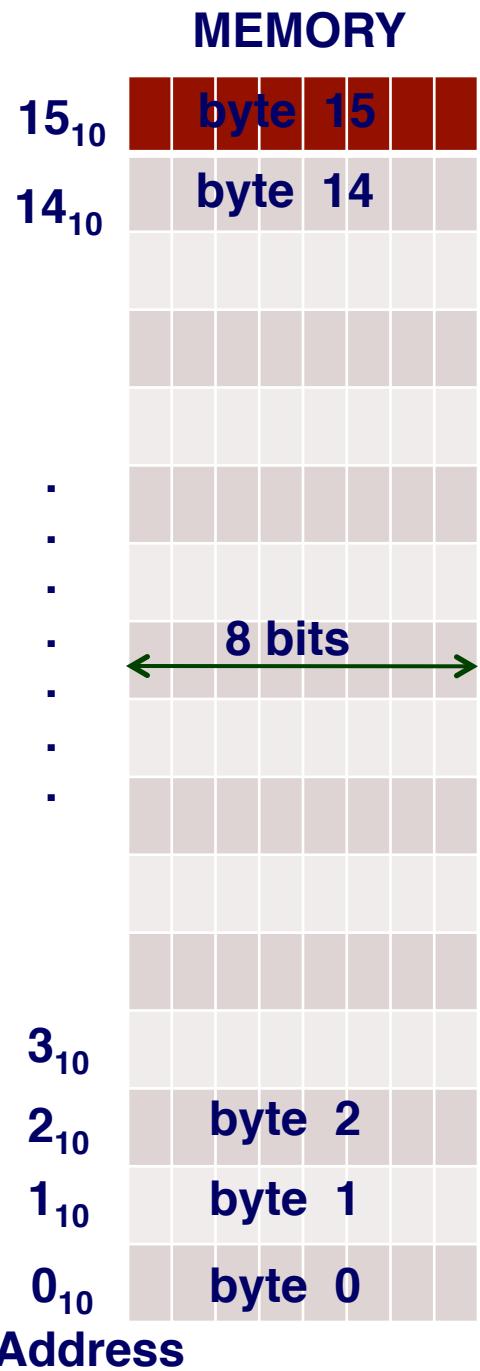
- Binary 00000000_2 to 11111111_2
- Decimal: 0_{10} to 255_{10}
- Hexadecimal 00_{16} to FF_{16}
- Memory addresses are referenced by their *byte* number
 - Example: a 16 byte memory
 - Memory address 15_{10}
 - = 16th byte (since memory starts at 0)
 - = binary address

$000000000000000000000000000000001111_2$ in bytes

= hex address $0x0000000F$ in bytes

This memory stores how many bits?

Answer: $8 * 16 = 128$ bits!



Machine Words

Each CPU Has “Word Size”

- Nominal size of integer-valued data
 - Including addresses
- Early microprocessors were 8-bit CPUs. Today, many embedded microcontrollers are still 8-bit.
- Intel evolved to 16-bit processors, and called these “words”.
 - `addw` and `movw` are assembly instructions where the “w” suffix stands for a 16-bit word and implies that 16-bit values are being added or moved
- Most machines became 32 bits (4 bytes)
 - `addl` and `movl` are assembly instructions where the “l” suffix stands for a “long word” and implies that 32-bit values are being added or moved
 - Limits addresses to 4GB
 - Becoming too small for memory-intensive applications

Machine Words

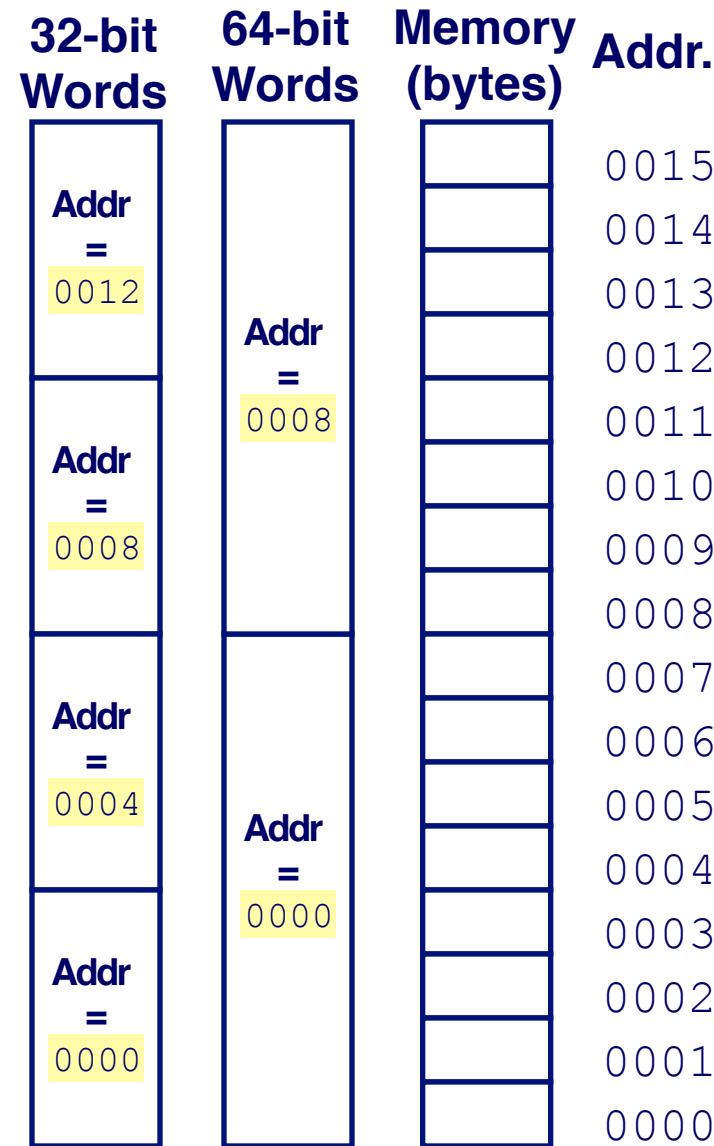
Each CPU Has “Word Size”

- Most newer systems are 64 bits (8 bytes)
 - Potentially address $\sim 1.8 \times 10^{19}$ bytes or 18 Exabytes
 - `addq` and `movq` are assembly instructions where the “q” suffix stands for “quad word” ($4 \times 16 = 64$) and implies that 64-bit values are being added or moved
- Machines support multiple data formats
 - Fractions or multiples of word size
 - Always integral number of bytes

Word-Oriented Memory Organization

Addresses Specify Byte Locations

- Address of first byte in word
- Addresses of successive words differ by 4 (32-bit) or 8 (64-bit)



Data Representations

Sizes of C Objects (in Bytes)

■ C Data Type	Intel_x64	Typical 32-bit	Intel IA32
● int	4	4	4
● long int	8	4	4
● char	1	1	1
● short	2	2	2
● float	4	4	4
● double	8	8	8
● long double	10/16	8	10/12
● char *	8	4	4

» Or any other pointer

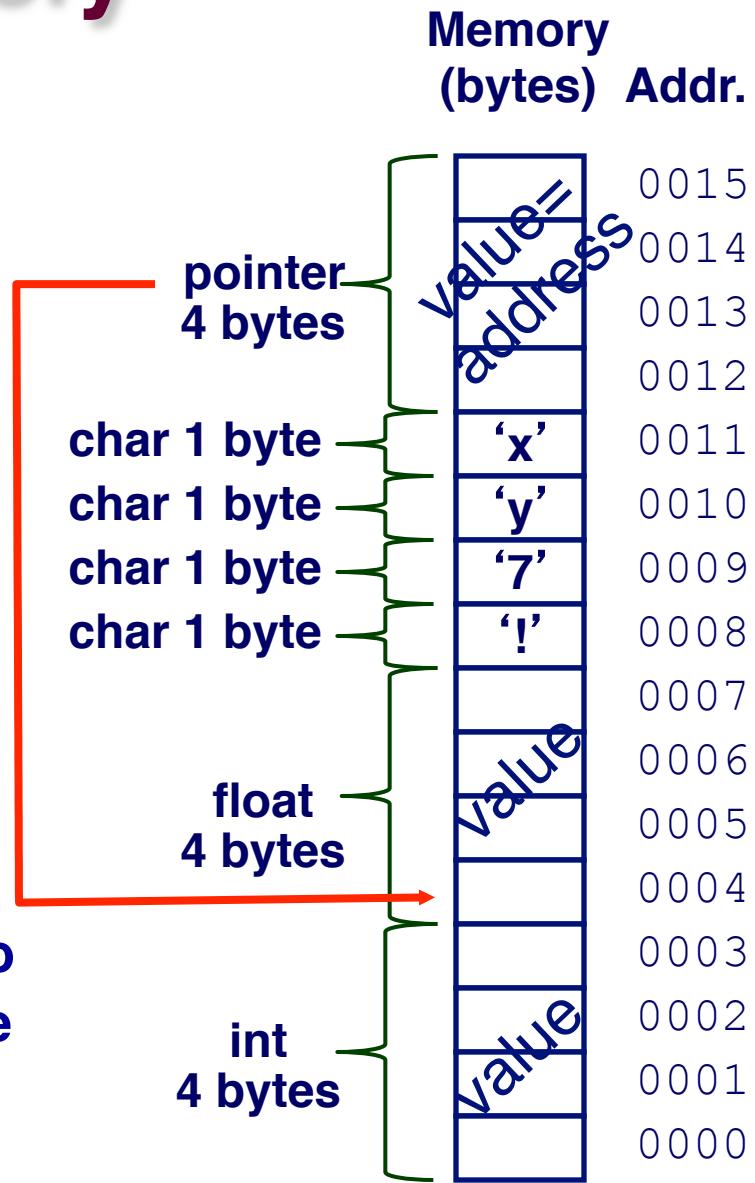
A float is a real # like -2.7143×10^{-3} . A double just provides more precision than a float, something like $2.71433273882 \times 10^{-3}$

A pointer stores an integer-like value that is interpreted as a memory location/address

Storing Data in Memory

IA32 Example:

- Address of int is 0x00000000
- Address of float is 0x00000004
- Address of character = '7' is 0x00000009
- Address of pointer is 0x0000000c
 - Note: the pointer points to another memory location, i.e. stores a memory location *address*
e.g. if pointer = 0x00000004 it means the pointer is pointing to the float! (actually the first byte of the float)



Pointers in C

To declare a pointer to say an integer, here's some code:

```
int x;  
int *p = &x;
```

int * p declares the variable p to be a pointer to an integer

note how it differs from the declaration ‘int p’ , which would just declare p to be an integer, not a pointer to an integer

&x means find the memory location/address of x

Thus, int *p = &x means initialize the value of the pointer p with the address of integer variable x.

Dereferencing Pointers in C

To find the value of a variable being pointed to, you must dereference a pointer. Here's some code:

```
int x;  
int *p = &x;  
int y = *p;
```

p is a pointer to an integer. To get the value of the integer that p is pointing to, you use the dereference operator (*).

note how * is used first to declare a pointer, and is used later to dereference a pointer to find the value of the variable it is pointing to

Here, y = x.

Pointer Values

Since the pointer value is an address, then it doesn't matter what type of variable is being pointed to, the size of the pointer value is always 4 bytes (IA32)

```
char c;  
int x;  
double d;  
char *p1 = &c;  
int *p2 = &x;  
double *p3 = &d;
```

p1 is a pointer to a one-byte long char, but is 4 bytes long, because the address of the char is 4 bytes long.

p3 is a pointer to an 8-byte long double, but is 4 bytes long, because the address of the double (lowest byte) is 4 bytes long.

Relations Between Logic Operations

DeMorgan's Laws

- Express $\&$ in terms of I , and vice-versa
 - $A \& B = \sim(\sim A \text{ I } \sim B)$
 - » A and B are true if and only if neither A nor B is false
 - $A \text{ I } B = \sim(\sim A \& \sim B)$
 - » A or B are true if and only if A and B are not both false

Exclusive-Or using Inclusive Or

- $A \wedge B = (\sim A \& B) \text{ I } (A \& \sim B)$
 - » Exactly one of A and B is true
 - » This is Shannon's circuit.
- $A \wedge B = (A \text{ I } B) \& \sim(A \& B)$
 - » Either A is true, or B is true, but not both

General Boolean Algebras

Operate on Bit Vectors

- Operations applied bitwise

$\&$	01101001	01101001	01010101
	01010101	$ $	01010101
	<hr/>	<hr/>	<hr/>
	01000001	0111101	10101010

All of the Properties of Boolean Algebra Apply

Using Boolean Operators for Representing & Manipulating Sets

Representation

- Width w bit vector represents subsets of $\{0, \dots, w-1\}$
- $a_j = 1$ if $j \in A$

01101001

$\{0, 3, 5, 6\}$

76543210

01010101

$\{0, 2, 4, 6\}$

76543210

Operations

- **&** Intersection 01000001 $\{0, 6\}$
- **|** Union 01111101 $\{0, 2, 3, 4, 5, 6\}$
- **^** Symmetric difference 00111100 $\{2, 3, 4, 5\}$
- **~** Complement 10101010 $\{1, 3, 5, 7\}$

Bit-Level Operations in C

Operations &, |, ~, ^ Available in C

- Apply to any “integral” data type
 - long, int, short, char
- View arguments as bit vectors
- Arguments applied bit-wise

- Examples (Char data type)

- $\sim 0x41 \rightarrow 0xBE$

$\sim 01000001_2 \rightarrow 10111110_2$

- $\sim 0x00 \rightarrow 0xFF$

$\sim 00000000_2 \rightarrow 11111111_2$

- $0x69 \& 0x55 \rightarrow 0x41$

$01101001_2 \& 01010101_2 \rightarrow 01000001_2$

- $0x69 | 0x55 \rightarrow 0x7D$

$01101001_2 | 01010101_2 \rightarrow 01111101_2$

Supplementary Slides

Integer Algebra vs. Boolean Algebra

Integer Arithmetic

- $\langle \mathbb{Z}, +, *, -, 0, 1 \rangle$ forms a “ring”
- Addition is “sum” operation
- Multiplication is “product” operation
- $-$ is additive inverse
- 0 is identity for sum
- 1 is identity for product

Boolean Algebra

- $\langle \{0,1\}, \text{I}, \&, \sim, 0, 1 \rangle$ forms a “Boolean algebra”
- OR is “sum” operation
- AND is “product” operation
- \sim is “complement” operation (not additive inverse)
- 0 is identity for sum
- 1 is identity for product

Boolean Algebra \approx Integer Ring

■ Commutativity

$$A \mid B = B \mid A$$

$$A \& B = B \& A$$

$$A + B = B + A$$

$$A * B = B * A$$

■ Associativity

$$(A \mid B) \mid C = A \mid (B \mid C)$$

$$(A \& B) \& C = A \& (B \& C)$$

$$(A + B) + C = A + (B + C)$$

$$(A * B) * C = A * (B * C)$$

■ Product distributes over sum

$$A \& (B \mid C) = (A \& B) \mid (A \& C)$$

$$A * (B + C) = A * B + A * C$$

■ Sum and product identities

$$A \mid 0 = A$$

$$A \& 1 = A$$

$$A + 0 = A$$

$$A * 1 = A$$

■ Zero is product annihilator

$$A \& 0 = 0$$

$$A * 0 = 0$$

■ Cancellation of negation

$$\sim(\sim A) = A$$

$$-(-A) = A$$

Boolean Algebra \neq Integer Ring

■ Boolean: *Sum distributes over product*

$$A \mid (B \& C) = (A \mid B) \& (A \mid C) \quad A + (B * C) \neq (A + B) * (A + C)$$

■ Boolean: *Idempotency*

$$A \mid A = A$$

$$A + A \neq A$$

- “A is true” or “A is true” = “A is true”

$$A \& A = A$$

$$A * A \neq A$$

■ Boolean: *Absorption*

$$A \mid (A \& B) = A$$

$$A + (A * B) \neq A$$

- “A is true” or “A is true and B is true” = “A is true”

$$A \& (A \mid B) = A$$

$$A * (A + B) \neq A$$

■ Boolean: *Laws of Complements*

$$A \mid \sim A = 1$$

$$A + \sim A \neq 1$$

- “A is true” or “A is false”

■ Ring: *Every element has additive inverse*

$$A \mid \sim A \neq 0$$

$$A + \sim A = 0$$

Boolean Ring

- $\langle \{0,1\}, \wedge, \&, I, 0, 1 \rangle$
- Identical to integers mod 2
- I is identity operation: $I(A) = A$
 $A \wedge A = 0$

Property

- Commutative sum
- Commutative product
- Associative sum
- Associative product
- Prod. over sum
- 0 is sum identity
- 1 is prod. identity
- 0 is product annihilator
- Additive inverse

Properties of & and \wedge

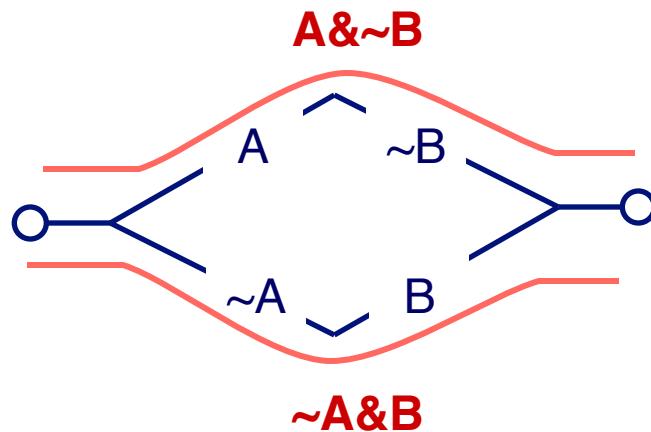
Boolean Ring

- $A \wedge B = B \wedge A$
- $A \& B = B \& A$
- $(A \wedge B) \wedge C = A \wedge (B \wedge C)$
- $(A \& B) \& C = A \& (B \& C)$
- $A \& (B \wedge C) = (A \& B) \wedge (A \& C)$
- $A \wedge 0 = A$
- $A \& 1 = A$
- $A \& 0 = 0$
- $A \wedge A = 0$

Application of Boolean Algebra

Applied to Digital Systems by Claude Shannon

- 1937 MIT Master's Thesis
- Reason about networks of relay switches
 - Encode closed switch as 1, open switch as 0



Connection when

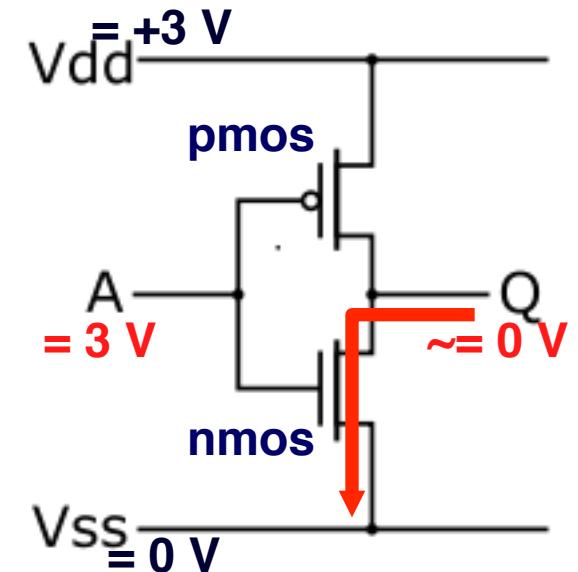
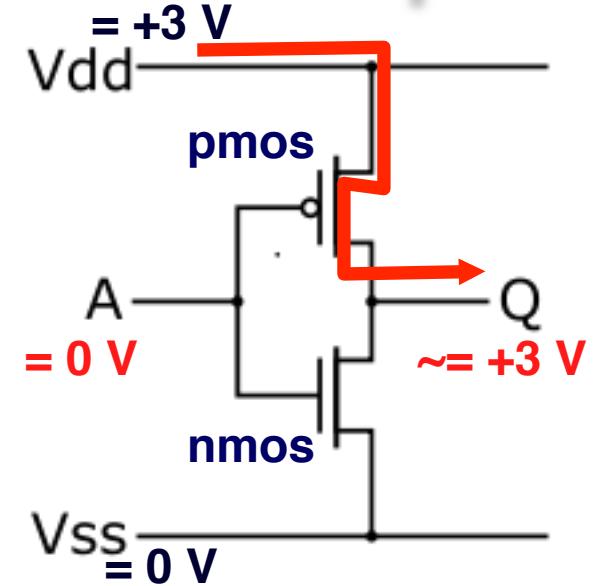
$$A \& \sim B \mid \sim A \& B$$

$$= A \wedge B$$

Binary Digital Logic Circuit Example

- **CMOS inverter circuit**

- **Vdd = Power = +3 Volts, Vss = Ground = 0 Volts**
- **Current flows from high voltage to low voltage**
- **When input A = 0 V (Ground), then nmos transistor is shut off, while the pmos transistor is turned on, so Q is connected to Vdd = +3 V**
- **When input A = +3 V (Power), then pmos is off and nmos is on, so Q is connected to ground and is pulled down to 0 V**
- **So Q is the opposite or inverse of A!**
i.e. $Q = \sim A$
- **AND, OR, NAND, NOR, ADD, SUBTRACT, etc. are also implemented this way...**



Byte-Oriented Memory Organization

Programs Refer to Virtual Addresses

- Conceptually very large array of bytes
- Actually implemented with hierarchy of different memory types
 - SRAM, DRAM, disk
 - Only allocate for regions actually used by program
- Address space private to particular “process”
 - Program being executed
 - Program can clobber its own data, but not that of others

Compiler + Run-Time System Control Allocation

- Where different program objects should be stored
- Multiple mechanisms: static, stack, and heap
- In any case, all allocation within single virtual address space